

Stop AI hallucinations before they reach your customers.
Score every response. **Catch what models won't.**

 THE PROBLEM

- LLMs fabricate citations, ISBNs, and study DOIs
- Models give medical dosages then add "I'm not a doctor"
- 87% of AI responses are too risky for direct delivery
- No audit trail = no compliance defense

 THE SOLUTION

- Score every AI response across 23 risk signals
- Auto-route unsafe responses to human reviewers
- Works with any model — internal, open-source, or cloud
- Full compliance audit trail on every response

23

RISK SIGNALS

25K

RESPONSES / SEC

<50ms

LATENCY

0

LLM DEPENDENCIES

— HOW IT WORKS



User Query

Customer asks your chatbot



LLM Response

Your model generates answer



Guardrail Scores

23 signals in <50ms



Safe? Deliver

Confident → show to user



Risky? Route

Flagged → human reviewer

— WHAT WE DETECT



Hallucinated facts & claims



Fabricated citations & ISBNs



Unsafe medical guidance



Risky financial advice



Unqualified legal opinions



Sycophantic agreement



Stale/outdated information



Disclaimer-then-advise pattern



Safety-critical content

INTEGRATION — 3 LINES OF CODE

```
// Your existing chatbot code — add 3 lines
const answer = await llm.chat(userQuestion);
const score = await fetch('https://api.guardrail.ai/check', { method: 'POST', body: JSON.stringify({ text: answer }) });
const { confidence, decision } = await score.json();
if (decision === 'deliver') send(answer); //  safe to show
if (decision === 'escalate') escalateToHuman(answer); //  needs review
```

— REAL AUDIT RESULTS — 5 LLMS, 100 RESPONSES

MODEL	CREATOR	PARAMS	SAFE TO DELIVER	FAILURE RATE
 Mistral	Mistral AI	7B	20%	80%
Phi-3	Microsoft	3.8B	15%	85%
Llama 3.2	Meta	3B	15%	85%
Gemma 2	Google	2B	10%	90%
Qwen 2.5	Alibaba	3B	5%	95%

Ready to secure your AI? **Let's talk.**

Live Playground →
symehmoo@gmail.com