



Guardrail AI

Real-time confidence scoring for AI outputs.
Stop hallucinations before they reach your
customers.

23 Risk
Signals

25K
Responses/sec

<50ms
Latency

Model-
Agnostic

symehmoo@gmail.com • guardrail-mvp-production.up.railway.app

GUARDRAIL AI

The Problem:

AI is confidently wrong.

- ✗ LLMs fabricate citations, ISBNs, and study DOIs
- ✗ Models give medical dosages, then add "I'm not a doctor"
- ✗ No safety layer between the AI and your customer
- ✗ A single bad response = lawsuit, lost trust, regulator attention

87%

OF AI RESPONSES WE
TESTED
WERE TOO RISKY TO
DELIVER

Based on 100 real responses
from 5 open-source LLMs

GUARDRAIL AI

The Solution:

Score every response. Catch what models won't.



Real-Time Scoring

Every AI response gets scored across 23 risk signals before it reaches your user. Under 50ms.



Smart Routing

Safe responses get delivered. Risky ones get routed to human reviewers. You set the thresholds.



Audit Trail

Every response scored, timestamped, and logged. Full compliance documentation for regulators.

💡 **Key differentiator:** Guardrail uses deterministic pattern matching — not another LLM. It's fast, it's predictable, and *it never hallucinates itself.*

GUARDRAIL AI

How It Works



Works with any LLM — OpenAI, Claude, Gemini, Llama, Mistral, or your fine-tuned internal model.

23 Signals We Detect



Hallucinated
facts & claims



Fabricated
citations &
ISBNs



Unsafe
medical
guidance



Risky
financial advice



Unqualified
legal opinions



Sycophantic
agreement



Stale /
outdated
information



Disclaimer-
then-advise
pattern



Safety-
critical content



Unverified
numerical
claims



AI identity
deflection



Hedging &
vague language

+ 11 more signals including knowledge cutoff detection, self-contradiction, overly brief responses, and formatting anomalies

GUARDRAIL AI

We Tested 5 LLMs.

87% of Responses Failed.

100 real responses • 20 adversarial prompts • 5 open-source models running locally via Ollama

100REAL LLM
RESPONSES**87%**FAILED
SAFETY
CHECK**100%**MEDICAL/LEGAL/FINANCIAL
FAILED

Model Safety Ranking

MODEL	CREATOR	PARAMETERS	SAFE TO DELIVER	FAILURE RATE
 Mistral	Mistral AI	7B	20%	80%
Phi-3	Microsoft	3.8B	15%	85%
Llama 3.2	Meta	3B	15%	85%
Gemma 2	Google	2B	10%	90%
Qwen 2.5	Alibaba	3B	5%	95%

Even the best model (Mistral 7B) failed 80% of safety checks. The "winner" is simply the least dangerous.

Key Findings



"I'm not a doctor" doesn't help

Every model added a medical disclaimer — then gave specific drug dosages anyway. Users act on the numbers, not the disclaimer.



4/5 models fabricated ISBNs

Asked for a real book's ISBN — at least 2 models generated fake numbers with high confidence. Users would unknowingly cite fake references.



More parameters ≠ safer

Google's 2B model beat Meta's 3B model on several prompts. Shorter responses = fewer unverified claims = higher safety score.



Verbose = dangerous

Models that gave longer responses scored lower. Every additional sentence is another opportunity to hallucinate.

GUARDRAIL AI

Integration: 3 Lines of Code

BEFORE — NO SAFETY LAYER

```
const answer =  
await  
llm.chat(question);  
res.send(answer);  
// 🚫 no check
```

AFTER — WITH GUARDRAIL

```
const answer = await  
llm.chat(question);  
const score = await  
guardrail.check(answer);  
  
if (score.decision ===  
'deliver')  
  res.send(answer); //  
  ✅ safe  
if (score.decision ===  
'escalate')  
  escalate(answer); //  
  🚫 human review
```

REST API — any
language

JavaScript
SDK

MCP Server (Claude,
Cursor, etc.)

Self-hosted
option

GUARDRAIL AI

Who Needs This



Healthcare & Pharma

AI patient triage, drug interaction checkers, symptom chatbots. Block unsafe dosage advice before it reaches patients.



Financial Services

Robo-advisors, loan chatbots, trading copilots. Prevent unqualified investment advice from reaching customers.



Legal Tech

Contract analysis, legal research assistants. Ensure AI outputs carry proper disclaimers and don't fabricate case law.



E-Commerce

Shopping assistants, product chatbots. Block hallucinated product specs and fake reviews from AI-generated content.



EdTech

AI tutors, homework helpers. Prevent factually wrong answers from reaching students at scale.



Enterprise SaaS

Any company adding "AI features" to their product. Add a safety net before shipping AI to customers.

What We've Built

Product

- Scoring API — deployed on Railway (production)
- JavaScript SDK — npm installable
- MCP Server — works with Claude, Cursor
- Interactive Playground — live demo
- Audit Toolkit — open-sourced on GitHub

Validation

- 100 real LLM responses audited
- 5 open-source models tested
- 23 detection signals implemented
- 25,000 req/sec throughput (benchmarked)
- Medium article & blog-ready report

Node.js **Railway** **REST API** **Heuristic**

Runtime

CI/CD

Integration

No LLM dependency

GUARDRAIL AI

Current Phase: **Hardening**

We're battle-testing the engine before scaling. Here's what we're actively working on.



Scoring Precision

Tuning 60+ heuristic weights using the "worst 15 responses" from our audit as adversarial benchmarks. Targeting <5% false positive rate.



Negation Logic

Upgrading claim extraction to distinguish "X is the capital" from "X is NOT the capital" — a common edge case in factual verification.



PostgreSQL Migration

Moving from JSON file storage to PostgreSQL on Railway for production-grade concurrency, persistence, and multi-tenant key management.



Larger Model Audits

Expanding from 2–7B models to 13B, 34B, and 70B parameter models. Testing 500+ prompts across 10+ models.



Enterprise Auth


Adding API key rotation, rate limiting, usage dashboards, and team management for enterprise customers.



Multi-language Support

Extending detection signals beyond English to support Spanish, French, German, Arabic, and Mandarin responses.

Product Roadmap



GUARDRAIL AI

Design Partner Program

We're selecting 3–5 companies to shape the product with us.

What You Get



Free access during the entire partner period (6 months)



Custom signals tuned to your specific domain and use case



Direct engineering access — weekly syncs, priority support



Custom audit report on your internal model's safety profile



Locked-in pricing — 50% discount on launch pricing, forever

What We Ask



Bi-weekly feedback — 30-min calls on what's working, what isn't



Real production traffic — run Guardrail on your actual AI pipeline



Case study permission — let us share (anonymized) results



Co-development — help prioritize the next detection signals



Ideal partners: Companies with live AI chatbots in healthcare, finance, legal, or customer support

GUARDRAIL AI

Business Model

STARTER	POPULAR	ENTERPRISE
Free forever	\$99 / month	Custom annual contract
✓ 1,000 checks / month ✓ 23 detection signals ✓ REST API access ✓ Community support ✗ Custom signals ✗ Analytics dashboard	✓ 50,000 checks / month ✓ 23 detection signals ✓ REST API + SDK ✓ Analytics dashboard ✓ Webhook alerts ✓ Priority support	✓ Unlimited checks ✓ Custom detection signals ✓ Self-hosted deployment ✓ SOC 2 compliance bundle ✓ SLA & dedicated support ✓ Team management & RBAC

Design partners receive 50% off Pro / Enterprise pricing for life.

GUARDRAIL AI

Why Guardrail AI

	GUARDRAIL AI	LLM-AS-JUDGE	MANUAL REVIEW	PROMPT ENGINEERING
Latency	✓ <50ms	✗ 2–10 sec	✗ Hours/days	~ Varies
Deterministic	✓ Always	✗ Non-deterministic	✓ Yes	✗ No
Cost per Check	✓ \$0.002	✗ \$0.01–0.05	✗ \$0.50–2.00	✓ Free
Audit Trail	✓ Full	~ Partial	✓ Full	✗ None
Scalability	✓ 25K/sec	~ Limited by API	✗ Not scalable	✓ Infinite
Can Hallucinate	✓ Never	✗ Yes — it's an LLM	✓ No	✗ Yes



The key insight: You can't trust an LLM to judge another LLM. Guardrail uses deterministic pattern matching — it's fast, cheap, and *never wrong about being wrong*.



Become a Design Partner. Shape the future of **AI** **safety.**

We're selecting 3–5 companies to build with.
Free access. Direct engineering support. Locked-in pricing
forever.

symehmoo@gmail.com



Live Playground → Try it now



Audit Report → GitHub

The question isn't whether your AI will make mistakes.
It's whether you'll catch them before your users do.